

# Mesh Painting on Subdivision Surfaces In Virtual Environments

\* Benjamin Gregorski    †Falko Kuester    \*Bernd Hamann    \*Kenneth I. Joy

\* Center for Image Processing and Integrated Computing (CIPIC)

Department of Computer Science, University of California, Davis, CA 95616-8562, USA

†Visualization and Interactive Systems Group

Department of Electrical and Computer Engineering, UC Irvine, CA 92697-2025

{gregorsk,hamann,joy}@cs.ucdavis.edu, fkuester@ieee.org

## Abstract

*Virtual environments allow us to model and design objects in a true 3D setting. We present an algorithm for interactive painting on subdivision surfaces within a semi-immersive virtual environment. A surface is first parameterized to create a set of texture maps that cover it. Polygons are then assigned texture coordinates and mapped into one of the textures. A set of line segments is used to represent a brush stroke across the surface. As the surface is painted, the texture map is updated to reflect its new color. Multiple texture maps covering the surface allow more detail to be painted in certain areas by using textures of different sizes. Within the virtual environment, the user is able to paint on the surface using a stylus tool that acts like a paint brush.*

**Keywords:** Virtual Reality, Immersive Environments, Interactive Modeling, Subdivision Surfaces, Mesh Painting

## 1. Introduction

Virtual environments (VEs) have been in use for many years in applications such as automotive and industrial design as well as advanced training and simulation. In these applications, VEs help to improve the efficiency of product design processes and aid in the learning process of new skills. Modeling applications, such as surface, editing that take advantage of VEs can greatly improve the speed and efficiency of the modeling process by adding an “extra dimension” making it easier to discern spatial relationships between objects and to interact with them.

Polygon meshes are a standard representation of 3D objects for applications in geometric modeling, animation, and grid generation. Sometimes, it is necessary to give models a realistic appearance instead of the standard look produced by smooth shaded polygons. Surface characteristics can be rendered using procedural shading and texture mapping methods to produce surfaces that simulate real-world material such as wood, marble, cloth, or metal. Another




Figure 1: Virtual Environment setup showing workbench, data gloves, and stylus

method for altering the appearance of a surface is to have an artist or designer “paint” these qualities onto the surface using a brush that applies colors or textures.

Subdivision surfaces have emerged as a powerful tool for surface modeling and editing. Lounsbery et al. [?] developed multiresolution analysis techniques using wavelets and used them for surface editing at coarse and fine levels. These ideas were also developed by Stollnitz et al. [?]. Modeling with subdivision surfaces was explored by Zorin et al. [?], where Loop surfaces [?] and Taubin smoothing [?] are the means to perform multi-resolution editing of a polygonal model. These techniques allow for interactive editing of subdivision surfaces on coarse and fine levels. Subdivision surfaces have also been extended to general design applications such as computer animation and geometric modeling. Hoppe et al. [?] developed rules for generating creases, darts, and sharp features that allow subdivisions surfaces to model a wider range of objects. Subdivision sur-

faces with sharp features were further developed by DeRose et al. [?], where they use subdivision techniques to generate models used in computer animation. Biermann et al. [?], have developed rules for controlling normal vectors on interior and boundary areas of surfaces.

In this paper, a method for painting on subdivision surfaces within virtual environments is presented. Our algorithm first determines a set of texture maps that cover the surface. Each polygon in the subdivision surface has texture coordinates that are associated with a texture map. Texture maps are referred to as *base textures*. When a point on the surface is painted, the texture coordinates of the point are determined, and the base texture corresponding to the polygon is modified to reflect the color change. As a user paints on a surface, the textures are updated and the painted model is rendered in real time. The system is integrated into a virtual environment using an immersive workbench and spatially tracked data gloves for interaction. A picture of our virtual environment configuration is shown in Figure 1. The application was developed as a plugin for the VirtualExplorer framework presented in [?].

## 2. Related Work

Mesh painting in 3D space is a common practice in companies that work on animation and special effects. Several commercial packages that allow users to paint on 3D objects are available. Hanrahan and Haeberli [?] have described a system for painting on 3D parameterized meshes using a 2D input device. Painting is performed directly on the mesh in a WYSIWYG (What-You-See-Is-What-You-Get) fashion. A user can manipulate the parameters used to shade a 3D object by applying pigments to its surface. The pigment has all the properties associated with material shading models such as diffuse and specular color and surface roughness. This idea was used by Agrawala et al. [?], along with a flood-fill algorithm for painting mesh vertices. Incremental drawing allows users to paint large meshes interactively without the use of expensive hardware. Since hardware has improved greatly in the past five years, their methods could probably be extended to larger meshes. In their system, a force feedback Polhemus device is used to paint on a triangle mesh.

Kuester et al. [?] have developed techniques for interactive modeling environments using immersive technologies. The *inTouch* system presented in [?] is a system for painting and modeling subdivision surfaces. This system uses a haptic device for multiresolution mesh editing and mesh painting. Ferley et al. [?] use isosurfaces to represent the surface being modeled. This approach enables a user to model surfaces of arbitrary shape and topology.




Figure 2: Subdivision of a quadrilateral and corresponding texture coordinates




Figure 3: Subdivision of a triangle and corresponding texture coordinates

## 3. Base Textures

Base textures as discussed earlier are a set of texture maps that cover a surface. The number and size of these texture maps determine how much detail can be painted onto a particular area of the surface. In order to determine how these base textures cover the surface, a surface parameterization needs to be defined. This parameterization maps a point on the surface to a texture coordinate in a specific base texture. For subdivision surfaces, the base mesh can be used as an initial parameterization. One base texture is associated with each face of the base mesh, and texture coordinates are associated on a per-face basis. As the mesh is subdivided, new texture coordinates are determined for the new faces, or children, by linearly subdividing the texture coordinates of the original faces or parents.

Subdivision of texture coordinates is shown for quadrilaterals in Figure 2 and for triangles in Figure 3. The left polygon shows the initial texture coordinates for the parent face, and the right polygon shows the texture coordinates of the child faces. It is important to note that the texture coordinates are assigned on a per-face and not a per-vertex basis even though this is not expressed in the figures. Texture coordinates in both parametric directions ( $u$  and  $v$ ) vary between 0 and 1.

All new faces that result from the subdivision process use




Figure 4: Base mesh, each base texture is a single color




Figure 5: Mesh from Figure 4 after 3 Catmull-Clark subdivisions

the same base texture as their “parent” face. This is illustrated in Figures 4 and 5 for a Catmull-Clark [?] subdivision surface. The base mesh of the surface, shown in Figure 4, has been colored such that no adjacent faces have the same color. Each face in the base mesh is assigned to a unique base texture map, and is assigned texture coordinates according to the scheme shown in Figure 2. The texture maps contain a single color that corresponds to the color for the face, and the faces in turn are colored using these texture maps.

As the mesh is subdivided, children inherit the base texture of their parent. The initial resolution of the base textures is 256x256 texels. The size of a base texture can be increased or decreased during the editing process by subsampling or supersampling the initial base texture. In either case, the texture coordinates of the faces associated with the base texture do not need to be changed.




Figure 6: Brush stroke and stroke segments

## 4. Painting Surfaces

### 4.1. Brush Strokes

To paint on a surface, the concept of a brush stroke similar to that presented by Gregory et al. [?] is used. A brush stroke is modeled by a series of line segments on the polygons of the subdivision surface, called *stroke segments*, and a *brush* that determines how the surface properties are modified. Each segment is confined to a single face in the mesh and represents a line in the face’s base texture. A segment is not a line between two points on the face but rather a line between two texture coordinates in the base texture associated with the face. This is similar to the approach used by Khodakovsky and Schröder [?] for editing fine underlying features on subdivision surfaces. However, instead of modifying geometry around a line segment, we modify surface attributes by editing the texture map around a line segment according to the different properties of a currently applied brush.

Figure 6 shows a brush stroke that spans several triangles. The brush stroke is broken into six segments marked A-F. Each segment is associated with one face. Each face is associated with one base texture and has separate texture coordinates for its vertices. A brush stroke is drawn by rasterizing each of its stroke segments in the corresponding texture maps and then modifying surrounding texels according to the brush properties.

When painting in 2D-space, movement vectors on the screen are translated into 3D movement vectors. This is achieved by mapping the start and end points of a mouse movement on the screen to points on the near plane of the camera. The process is illustrated in Figure 7. Two rays from the camera pass through the points on the near plane and are intersected with the object being painted. This yields two points in physical object space that define the movement vector in 3D space.

The movement vector is projected back onto the surface to form the line segments of the brush stroke. Beginning with the face intersected by the starting ray, the movement vector is projected into the plane defined by the face to form the stroke segment for that face. The next face is found by determining which edge or vertex of the face the projected




Figure 7: Creation of 3D movement vector

vector intersects. The face on the opposite side of the edge or vertex becomes the next face. The new movement vector is projected onto this new face, and the process continues until the length of the movement vector is below a threshold.

## 4.2. Brush Characteristics

The rasterization of a stroke segment in a texture map touches a set of texels called the *stroke texels*. Figure 8 shows a single stroke segment in a triangle and the corresponding stroke texels in the texture map. The stroke segment is shown in physical and texture space along with the mapping of texture coordinates from physical space to texture space. The properties of a brush define how the stroke texels and the surrounding texels are modified. These are the properties of a brush:

1. **Shape and size.** The shape and size of a brush define which surrounding texels are modified. The size of a brush is given in physical space, and is described by a bounding sphere. All texels that lie within this bounding sphere can be modified by the brush. The shape of the brush determines which texels within its bounding sphere are modified. A circular brush, for example, selects which texels are modified based on the distance between the center of the texel and the center of the corresponding stroke texel.
2. **Color function.** The color function defines the color of the brush at a particular position within its extent, i.e., the volume of space defined by its size and shape. The color function can be a basic function, such as a checker or stripe pattern, or a more complex function such as procedural marble or wood.
3. **Blend function.** The blend function defines how the brush attributes, as defined by the color function, are combined with the existing texel. The blend function can either replace the existing texel attributes or blend them with the brush attributes.




Figure 8: Brush stroke and stroke texels

Brushes modify the diffuse color of the surface, represented by an RGBA tuple. It is also possible to use the brushes to paint different material properties, such as specular highlights and bump map information, onto the surface. Interactive rendering can be achieved by using multi-texturing and normal mapping to render the surface.

## 4.3. Texel Modification

The texels surrounding the stroke texels are modified according to the characteristics of the brush stroke. The surrounding texels fall into one of three categories:

1. Texels that lie in the same face and same texture as the stroke texel.
2. Texels that lie in a different face and same texture as the stroke texel.
3. Texels that lie in a different face and different texture as the stroke texel.

These texels are found by placing the brush-specific bounding sphere at each stroke texel and using the size of the brush to determine the surface faces that are touched. The position of the brush's bounding sphere in physical space is found by mapping the position of the stroke texel from texture to physical space. This is done as follows:

1. The texel's indices into the texture map are converted to texture coordinates that lie in the interval  $[0,1]$ .
2. The barycentric coordinates of the texel's texture coordinates are determined using the texture coordinates of the face's vertices.
3. The barycentric coordinates are used to compute the texel's position in physical space from the vertices of the face. The position is computed for the center of the texel.




Figure 9: Faces on the far side are not painted when a breadth-first search is used.




Figure 10: Both sides of the surface (dashed) are painted.

Starting with the initial face, the surrounding faces in the mesh are searched to find those faces that are partially or completely inside the sphere. The search is performed in a breadth-first fashion. This method prevents faces from being painted incorrectly when two parts of the surface are close together. This is illustrated in Figure 9: The sphere created at the stroke texel contains faces from the far side of the surface. The breadth-first search starting at the initial face will not include the faces on the far side of the surface. As a result, these faces will not be incorrectly painted when the brush is moved over the surface. A special case is shown in Figure 10. The brush contains all of the faces on some path from the painted side of the surface to the far side of the surface. In this case, the breadth-first search used to find the painted faces will find those faces on the far side of the surface correctly.

## 5. Virtual Reality Modeling

### 5.1. Overview

We have integrated our system into a Virtual Reality environment using an immersive workbench from Fakespace Corporation. The semi-immersive environment allows stereo viewing using headtracked shutter glasses. This has many advantages over traditional editing environments where the user must interact with objects projected onto a 2D screen. In our system, a user can interact with the environment using spatially tracked data gloves and a sty-




Figure 11: Decomposition of painting gesture into surface and brush vectors

lus. The interface was implemented using the *VirtualExplorer* framework developed by researchers at UC Davis and Irvine. The VirtualExplorer framework is an object-oriented, customizable, plugin-based framework for VR applications. The basis of VirtualExplorer is a run-time plugin system that allows users to dynamically load, unload, enable, or disable different modules of functionality. Our algorithms are implemented as a plugin for the VirtualExplorer framework.

### 5.2. Painting Gestures

The painting process is performed as the user moves the brush through 3D space. The brush is drawn at the position of the current stroke texel being painted to indicate where the brush is during painting.

Painting gestures are broken down into gestures that move the brush along the surface and gestures that move the brush towards or away from the surface. Gestures that move the brush along the surface are used to determine the stroke segments that compose the current brush stroke. The movement vector associated with this gesture is called the *surface vector*. Gestures that move the brush towards or away from the surface determine how much of the brush is in contact with the surface, and thus the section of the surface that is painted. The movement vector associated with this gesture is called the *brush vector*.

The decomposition of a movement in 3D space into the surface vector and brush vector is illustrated in Figure 11. A movement vector is initiated at some point on the surface. This movement vector is divided into components tangent and perpendicular to the surface. The component along the surface is found by the same algorithm for determining the stroke segments as described in Section 4.1.

### 5.3. Applying Paint

Painting gestures are used to determine which portion of a surface is painted and how much paint is applied to this area. As the brush is moved across the surface in a painting gesture, the individual components of the gesture can be used to determine the extent to which a particular brush interacts with the surface. The brush vector is used to determine how much of the brush is touching the surface, and thus how many texels are painted. The surface vector is used to determine the stroke segments.

### 5.4. Data Gloves

The data gloves allow the user to interact with the virtual environment in a two-handed free-form manner. The spatially tracked data gloves provide position information for the hand and pinch information for the fingers. The data gloves are used for these purposes:

1. Interactive viewpoint selection
2. Selection of a surface or surface feature for editing

Interactive viewpoint selection is performed by either translating the viewpoint through the scene or rotating the viewpoint around a virtual trackball located at some point in the environment. Viewpoint translation is accomplished by pinching the thumb and index fingers of the data glove together and moving the glove in space. Viewpoint rotation is accomplished by pinching together the thumb and index fingers of both hands and moving them in space. This motion determines a rotation axis and an angle of rotation. In scenes with multiple objects, the data gloves are used to select which surface to edit.

### 5.5. Stylus

The stylus tool allows a user to select objects in the scene using a method similar to pointing out objects with a laser pointer. This enables a user to select distant objects, to bring distant objects into the foreground, or to move them to specific locations. The stylus provides the direction that it is pointing towards and pinch information for its button. Direction information is available as a quaternion rotation and as a tuple consisting of direction, pitch, and roll. The stylus is used for these purposes:

1. Selecting objects for modeling
2. Painting objects using the stylus as a paint brush

The stylus is represented as a ray in 3D space, i.e., it has a 3D position and direction. Object painting is performed by pointing the stylus at the surface to select a starting point and then moving the stylus's ray along the surface to trace

out the path of the brush. The intersection of the stylus's ray with the surface is used to create a movement vector in 3D space. This 3D movement vector is painted on the surface in the manner described in Section 4.1.

Painting using the stylus can be done in either *laser-pointer* mode or *paint-brush* mode. In laser-pointer mode, the position of the stylus is not taken into account, and only the stylus's direction is used to construct the brush strokes on the surface. This allows the user to make broader strokes over the surface of the object. In paint-brush mode, the position of the stylus is used to determine how the brush interacts with the surface. In this painting style, the brush has the notion of how much of it is on the surface, and thus what portion of the surface is painted. This allows the user to paint finer strokes over the object. The stylus tool enables a user to select the points in 3D space that describe the movement vector within the space of the object rather than on a 2D projection of the object.

## 6. Results

Figure 12 shows a painted bunny. The bunny model was obtained from Stanford University's model archive and simplified using Michael Garland's qslim software. The base mesh consists of 500 triangles, and it was subdivided three times using Loop's subdivision method. There are 500 base textures each with 64x64 texels. The base textures are assigned colors such that as few adjacent faces as possible share the same color. The letters on the bunny were painted with a spherical brush. The blend function of the brush replaces the existing texel's color with the brush color. Figure 13 shows the bicycle seat model from Figure 4 subdivided three times using Catmull-Clark subdivision. There are 18 base textures with 256x256 texels. The small strokes were painted with a brush of half the size used for the big strokes.

## 7. Conclusions and Future Work

We have presented a method for interactive painting on a subdivision surface in a virtual environment. Our algorithm starts by parameterizing the surface using the base mesh and forming a set of textures that cover the surface. The user is able to interactively paint on the surface by modifying the underlying texture maps. The VR environment provides a user with a much better spatial perception of the model and allows one to paint surfaces in 3D space.

Future work will be directed at the development of new interaction paradigms in VEs and a user study of these paradigms. New modeling algorithms that take maximal advantage of a VR environment need to be investigated, specifically with regard to the advantages and limitations of our current hardware setup. A user interface for the VE that





Figure 12: Painted bunny. Brush shape = spherical, Blend function = replace




Figure 13: Painted bike seat. Brush shape = spherical, Blend function = replace

provides the user access to the full range of tools available to an artist has to be developed. This includes the addition of painting specific items such color selection, brush selection, and a virtual palette. The development of a richer user interface will greatly add to the capabilities of the system.

## Acknowledgments

This work was supported by the National Science Foundation under contracts ACI 9624034 (CAREER Award), through the Large Scientific and Software Data Set Visualization (LSSDSV) program under contract ACI 9982251, and through the National Partnership for Advanced Computational Infrastructure (NPACI); the Office of Naval Research under contract N00014-97-1-0222; the Army Research Office under contract ARO 36598-MA-RIP; the NASA Ames Research Center through an NRA award under contract NAG2-1216; the Lawrence Livermore National Laboratory under ASCI ASAP Level-2 Memorandum Agreement B347878 and under Memorandum Agreement B503159; the Lawrence Berkeley National Laboratory; the Los Alamos National Laboratory; and the North Atlantic Treaty Organization (NATO) under contract CRG.971628. We also acknowledge the support of ALSTOM Schilling Robotics and SGI. We thank the members of the Visualization and Graphics Research Group at the Center for Image Processing and Integrated Computing (CIPIC) at the University of California, Davis.

## References

- [1] Maneesh Agrawala, Andrew C. Beers, and Marc Levoy. 3D painting on scanned surfaces. In Pat Hanrahan and Jim Winget, editors, *1995 Symposium on Interactive 3D Graphics*, pages 145–150. ACM SIGGRAPH, April 1995.
- [2] Henning Biermann, Adi Levin, and Denis Zorin. Piecewise smooth subdivision surfaces with normal control. In Kurt Akeley, editor, *Siggraph 2000, Computer Graphics Proceedings*, pages 113–120. ACM Press / ACM SIGGRAPH / Addison Wesley Longman, 2000.
- [3] Edwin Catmull and Jim Clark. Recursively generated B-spline surfaces on arbitrary topological meshes. *Computer-Aided Design*, 10, September 1978.
- [4] Tony DeRose, Michael Kass, and Tien Truong. Subdivision surfaces in character animation. In Michael Cohen, editor, *SIGGRAPH 98 Conference Proceedings*. ACM SIGGRAPH, Addison Wesley, 1998.

- [5] Eric Ferley, Marie-Paule Cani, and Jean-Dominique Gascuel. Practical volumetric sculpting. In *Proceedings of Implicit Surface '99*, Sep 1999.
- [6] Arthur Gregory, Stephen Ehmann, and Ming Lin. in-Touch: Interactive multiresolution modeling and 3D painting with a haptic interface. In Steven Feiner and Daniel Thalmann, editors, *Proceedings of the 2000 IEEE Conference on Virtual Reality (VR-00)*. IEEE.
- [7] Pat Hanrahan and Paul Haeberli. Direct WYSIWYG painting and texturing on 3D shapes. In Forest Baskett, editor, *Computer Graphics (SIGGRAPH '90 Proceedings)*, volume 24, August 1990.
- [8] Hugues Hoppe, Tony DeRose, Tom Duchamp, Mark Halstead, Hubert Jin, John McDonald, Jean Schweitzer, and Werner Stuetzle. Piecewise smooth surface reconstruction. In Andrew Glassner, editor, *Proceedings of SIGGRAPH '94 (Orlando, Florida, July 24–29, 1994)*. ACM SIGGRAPH, ACM Press.
- [9] Andrei Khodakovsky and Peter Schröder. Fine level feature editing for subdivision surfaces. In Willem F. Bronsvort and David C. Anderson, editors, *Proceedings of the Fifth Symposium on Solid Modeling and Applications (SSMA-99)*. ACM Press.
- [10] Falko Kuester, Mark A. Duchaineau, Bernd Hamann, Kenneth I. Joy, and Kwan-Liu Ma. The designer-workbench: towards real-time immersive modeling. In J. O. Merritt, S. A. Benton, A. J. Woods, and M. T. Bolas, editors, *Stereoscopic Displays and Virtual Reality Systems VII*, volume SPIE Vol. 3957. The International Society for Optical Engineering, 2000.
- [11] Falko Kuester, Bernd Hamann, and Kenneth I. Joy. Virtualexplorer: A plugin-based virtual reality framework. In R.F. Erbacher, P.C. Chen, M. Groehn, J.C. Roberts, and C.M. Wittenbrink, editors, *Proceedings of SPIE*, San Jose, California, USA, 2001. SPIE - The International Society of Optical Engineering.
- [12] Charles Loop. Smooth subdivision surfaces based on triangles. Master's thesis, 1987. University of Utah, Department of Mathematics.
- [13] Michael Lounsbery, Tony D. DeRose, and Joe Warren. Multiresolution analysis for surfaces of arbitrary topological type. *ACM Transactions on Graphics*, 16(1):34–73, January 1997.
- [14] Eric J. Stollnitz, Tony D. DeRose, and David H. Salesin. *Wavelets for Computer Graphics: Theory and Applications*. Morgan Kaufmann, San Francisco, CA, 1996.
- [15] Gabriel Taubin. A signal processing approach to fair surface design. In Robert Cook, editor, *SIGGRAPH 95 Conference Proceedings*, Annual Conference Series, pages 351–358. ACM SIGGRAPH, Addison Wesley, August 1995.
- [16] Denis Zorin, Peter Schröder, and Wim Sweldens. Interactive multiresolution mesh editing. In Turner Whitted, editor, *SIGGRAPH 97 Conference Proceedings*, pages 259–268. ACM SIGGRAPH, Addison Wesley, August 1997.